

## Week 8.2 Assignment

- Joshua Burden
- Bellevue University
- DSC550 Data Mining
- Dr. Brett Werner
- 10/23/2022

**Begin Milestone 1 with a 250-500-word narrative describing your original idea for the analysis/model building business problem.**

**Clearly identify the problem you will address and the target for your model.**

### **Background:**

During the pandemic an increase in the need for health care professionals was required. The dataset collected is a modified synthetic dataset from IBM's Watson to show a useful insight into the attrition rate for healthcare workers.

### **Problem:**

The data set includes information about the attrition rate for employees within the healthcare field. The meaning of employee attrition is the departure of employees from the organization for any reason whether that be voluntary or involuntary, including resignation, termination, death, or retirement. Companies to avoid attrition rates being too high is to replace those who are either leaving voluntarily or involuntary. The data set should provide insights into whether a company in the healthcare field was replacing their employees that were leaving the field, or if they continued to have a gradual but deliberate reduction in staff for any reason.

### **Original Idea:**

The idea behind this data set is to discover whether certain roles within the healthcare industry, hours worked, age of an employee, or any other qualifying data points stand out as to why the healthcare industry had any determining factor on whether a person was to leave their field, while also predicting whether the employee was eventually replaced.

### **Dataset:**

This dataset contains employee and company data useful for supervised ML, unsupervised ML, and analytics. Attrition - whether an employee left or not - is included and can be used as the

target variable. The data is synthetic and based on the IBM Watson dataset for attrition. Employee roles and departments were changed to reflect the healthcare domain. Also, known outcomes for some employees were changed to help increase the performance of ML models

Then, do a graphical analysis creating a minimum of four graphs.

Label your graphs appropriately and explain/analyze the information provided by each graph.

```
In [ ]: import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
import numpy as np
import plotly.express as px
```

```
In [ ]: data_df = pd.read_csv('./DATA/watson_healthcare_modified.csv')
data_df.head()
```

```
Out[ ]:
```

	EmployeeID	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educational
0	1313919	41	No	Travel_Rarely	1102	Cardiology	1	
1	1200302	49	No	Travel_Frequently	279	Maternity	8	
2	1060315	37	Yes	Travel_Rarely	1373	Maternity	2	
3	1272912	33	No	Travel_Frequently	1392	Maternity	3	
4	1414939	27	No	Travel_Rarely	591	Maternity	2	

5 rows × 35 columns

```
In [ ]: print("Number of duplicated data: "+str(data_df.duplicated().sum()))
```

Number of duplicated data: 0

```
In [ ]: data_df.isnull().sum()
```

```
Out[ ]: EmployeeID      0
        Age            0
        Attrition      0
        BusinessTravel 0
        DailyRate      0
        Department     0
        DistanceFromHome 0
        Education      0
        EducationField  0
        EmployeeCount   0
        EnvironmentSatisfaction 0
        Gender         0
        HourlyRate      0
        JobInvolvement  0
        JobLevel        0
        JobRole         0
        JobSatisfaction 0
        MaritalStatus   0
        MonthlyIncome   0
        MonthlyRate     0
        NumCompaniesWorked 0
        Over18          0
        OverTime        0
        PercentSalaryHike 0
        PerformanceRating 0
        RelationshipSatisfaction 0
        StandardHours   0
        Shift           0
        TotalWorkingYears 0
        TrainingTimesLastYear 0
        WorkLifeBalance 0
        YearsAtCompany  0
        YearsInCurrentRole 0
        YearsSinceLastPromotion 0
        YearsWithCurrManager 0
        dtype: int64
```


```
In [ ]: data_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1676 entries, 0 to 1675
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   EmployeeID                           1676 non-null   int64
1   Age                                   1676 non-null   int64
2   Attrition                            1676 non-null   object
3   BusinessTravel                       1676 non-null   object
4   DailyRate                            1676 non-null   int64
5   Department                           1676 non-null   object
6   DistanceFromHome                     1676 non-null   int64
7   Education                             1676 non-null   int64
8   EducationField                       1676 non-null   object
9   EmployeeCount                        1676 non-null   int64
10  EnvironmentSatisfaction               1676 non-null   int64
11  Gender                               1676 non-null   object
12  HourlyRate                           1676 non-null   int64
13  JobInvolvement                       1676 non-null   int64
14  JobLevel                             1676 non-null   int64
15  JobRole                              1676 non-null   object
16  JobSatisfaction                      1676 non-null   int64
17  MaritalStatus                       1676 non-null   object
18  MonthlyIncome                       1676 non-null   int64
19  MonthlyRate                          1676 non-null   int64
20  NumCompaniesWorked                  1676 non-null   int64
21  Over18                              1676 non-null   object
22  OverTime                             1676 non-null   object
23  PercentSalaryHike                   1676 non-null   int64
24  PerformanceRating                   1676 non-null   int64
25  RelationshipSatisfaction              1676 non-null   int64
26  StandardHours                       1676 non-null   int64
27  Shift                               1676 non-null   int64
28  TotalWorkingYears                   1676 non-null   int64
29  TrainingTimesLastYear                1676 non-null   int64
30  WorkLifeBalance                     1676 non-null   int64
31  YearsAtCompany                      1676 non-null   int64
32  YearsInCurrentRole                   1676 non-null   int64
33  YearsSinceLastPromotion              1676 non-null   int64
34  YearsWithCurrManager                 1676 non-null   int64
dtypes: int64(26), object(9)
memory usage: 458.4+ KB
```

## Visualization 1

```
In [ ]: plt.figure(figsize=(20,10))
sns.countplot(x='Age',hue='Attrition',data=data_df)
```

```
Out[ ]: <AxesSubplot:xlabel='Age', ylabel='count'>
```



## Visualization 2


```
In [ ]: data_df.groupby('Attrition')['MonthlyIncome'].mean().sort_values().reset_index()
```

Out[ ]:

	Attrition	MonthlyIncome
0	Yes	4024.246231
1	No	6852.301963

```
In [ ]: plt.figure(figsize=(20,10))
sns.countplot(x='DistanceFromHome',hue='Attrition',data=data_df)
```

Out[ ]: <AxesSubplot:xlabel='DistanceFromHome', ylabel='count'>



## Visualization 3

```
In [ ]: px.histogram(data_df,x="Department",color="Attrition",barmode="group",text_auto=".2f",  
                    title = "Percentage of Department Type")
```


```
In [ ]: px.histogram(data_df,x="EducationField",color="Attrition",barmode="group",text_auto=".  
                    title = "Percentage of EducationField Type")
```

```
In [ ]: px.histogram(data_df,x="JobRole",color="Attrition",barmode="group",text_auto=".2f",ten  
                    title = "Percentage of EducationField Type")
```

## Visualization 4

```
In [ ]: plt.figure(figsize=(10,10))  
sns.countplot(x='Gender',hue='Attrition',data=data_df)
```

```
Out[ ]: <AxesSubplot:xlabel='Gender', ylabel='count'>
```



### Breakdown of all the available datapoints

```
In [ ]: plt.figure(figsize=(30,50))
        for index,column in enumerate(num_col):
            plt.subplot(5,5,index+1)
            sns.countplot(data=num_col,x=column)
            plt.xticks(rotation = 90)
        plt.tight_layout(pad = 1.0)
        plt.show()
```

```
-----
NameError                                Traceback (most recent call last)
c:\Users\Joshu\Desktop\Masters\DSC550\JoshuaBurdenAssignment8-2.ipynb Cell 25 in <cel
1 line: 2>()
      1 <a href='vscode-notebook-cell:/c%3A/Users/Joshu/Desktop/Masters/DSC550/JoshuaBu
      2 rdenAssignment8-2.ipynb#X33sZm1sZQ%3D%3D?line=0'>1</a> plt.figure(figsize=(30,50))
----> <a href='vscode-notebook-cell:/c%3A/Users/Joshu/Desktop/Masters/DSC550/JoshuaBu
      3 rdenAssignment8-2.ipynb#X33sZm1sZQ%3D%3D?line=1'>2</a> for index,column in enumerate
      4 (num_col):
      5     <a href='vscode-notebook-cell:/c%3A/Users/Joshu/Desktop/Masters/DSC550/JoshuaBu
      6 rdenAssignment8-2.ipynb#X33sZm1sZQ%3D%3D?line=2'>3</a>         plt.subplot(5,5,index+1)
      7     <a href='vscode-notebook-cell:/c%3A/Users/Joshu/Desktop/Masters/DSC550/JoshuaBu
      8 rdenAssignment8-2.ipynb#X33sZm1sZQ%3D%3D?line=3'>4</a>         sns.countplot(data=num_co
      9 l,x=column)

NameError: name 'num_col' is not defined
<Figure size 2160x3600 with 0 Axes>
```

## Observations:

- Maternity departments had the highest rate of attrition followed by cardiology and neurology
- attrition rates had the highest peak at 29 years old
- 26-35 years old saw the highest range of attrition
- 42 years old and older saw the least attrition rates
- More men were likely to leave than women but Men also were more accounted for than women in the healthcare field
- Human resources were the least likely to have people quit
- Life Sciences were the Education field with the highest amount of attrition
- people that lived closer to their jobs were more likely to leave

## Milestone 2

```
In [ ]: data_df.head()
```

Out [ ]:

	EmployeeID	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education
0	1313919	41	No	Travel_Rarely	1102	Cardiology	1	
1	1200302	49	No	Travel_Frequently	279	Maternity	8	
2	1060315	37	Yes	Travel_Rarely	1373	Maternity	2	
3	1272912	33	No	Travel_Frequently	1392	Maternity	3	
4	1414939	27	No	Travel_Rarely	591	Maternity	2	

5 rows × 35 columns

In [ ]: `data_df.shape`

Out [ ]: (1676, 35)

Dropping data columns that don't provide much value or context to the data

In [ ]: `#drop some of the values`  
`data_df1 = data_df.drop(['EmployeeID', 'Over18', 'EmployeeCount', 'StandardHours'], axis=1)`  
`data_df1.head()`

Out [ ]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Education
0	41	No	Travel_Rarely	1102	Cardiology	1	2	Life Sci
1	49	No	Travel_Frequently	279	Maternity	8	1	Life Sci
2	37	Yes	Travel_Rarely	1373	Maternity	2	2	
3	33	No	Travel_Frequently	1392	Maternity	3	4	Life Sci
4	27	No	Travel_Rarely	591	Maternity	2	1	Me

5 rows × 31 columns

In [ ]: `#change attrition rates from no/yes to 0/1`  
`data_df1['Attrition'] = data_df1['Attrition'].str.replace('Yes', str(1))`  
`data_df1['Attrition'] = data_df1['Attrition'].str.replace('No', str(0))`  
`data_df1['Attrition'] = data_df1['Attrition'].astype('int')`

converted attritions yes/no values and replaced them with 0/1 and set type to int

Look at shape and values of columns


In [ ]: `data_df1.head()`



Out[ ]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Education
0	41	0	Travel_Rarely	1102	Cardiology	1	2	Life Sci
1	49	0	Travel_Frequently	279	Maternity	8	1	Life Sci
2	37	1	Travel_Rarely	1373	Maternity	2	2	
3	33	0	Travel_Frequently	1392	Maternity	3	4	Life Sci
4	27	0	Travel_Rarely	591	Maternity	2	1	Me

5 rows × 31 columns



In [ ]:

```
data_df1.describe().T
```

Out[ ]:

	count	mean	std	min	25%	50%	75%	max
Age	1676.0	36.866348	9.129126	18.0	30.00	36.0	43.00	60.0
Attrition	1676.0	0.118735	0.323573	0.0	0.00	0.0	0.00	1.0
DailyRate	1676.0	800.557876	401.594438	102.0	465.00	796.5	1157.00	1493.0
DistanceFromHome	1676.0	9.221957	8.158118	1.0	2.00	7.0	14.00	29.0
Education	1676.0	2.907518	1.025835	1.0	2.00	3.0	4.00	5.0
EnvironmentSatisfaction	1676.0	2.714797	1.097534	1.0	2.00	3.0	4.00	4.0
HourlyRate	1676.0	65.470167	20.207572	30.0	48.00	65.5	83.00	100.0
JobInvolvement	1676.0	2.724940	0.714121	1.0	2.00	3.0	3.00	4.0
JobLevel	1676.0	2.066826	1.113423	1.0	1.00	2.0	3.00	5.0
JobSatisfaction	1676.0	2.738663	1.104005	1.0	2.00	3.0	4.00	4.0
MonthlyIncome	1676.0	6516.512530	4728.456618	1009.0	2928.25	4899.0	8380.25	19999.0
MonthlyRate	1676.0	14287.019690	7138.857783	2094.0	7993.00	14269.5	20462.00	26999.0
NumCompaniesWorked	1676.0	2.662291	2.477704	0.0	1.00	2.0	4.00	9.0
PercentSalaryHike	1676.0	15.196897	3.646550	11.0	12.00	14.0	18.00	25.0
PerformanceRating	1676.0	3.150358	0.357529	3.0	3.00	3.0	3.00	4.0
RelationshipSatisfaction	1676.0	2.718377	1.078162	1.0	2.00	3.0	4.00	4.0
Shift	1676.0	0.806086	0.855527	0.0	0.00	1.0	1.00	3.0
TotalWorkingYears	1676.0	11.338902	7.834996	0.0	6.00	10.0	15.00	40.0
TrainingTimesLastYear	1676.0	2.805489	1.288431	0.0	2.00	3.0	3.00	6.0
WorkLifeBalance	1676.0	2.766110	0.702369	1.0	2.00	3.0	3.00	4.0
YearsAtCompany	1676.0	7.033413	6.098991	0.0	3.00	5.0	10.00	40.0
YearsInCurrentRole	1676.0	4.264916	3.627456	0.0	2.00	3.0	7.00	18.0
YearsSinceLastPromotion	1676.0	2.200477	3.229587	0.0	0.00	1.0	3.00	15.0
YearsWithCurrManager	1676.0	4.135442	3.559662	0.0	2.00	3.0	7.00	17.0

check correlation of data columns

In [ ]:

data\_df1.corr()

Out[ ]:

	Age	Attrition	DailyRate	DistanceFromHome	Education	Environment
Age	1.000000	-0.239984	0.001441	-0.010079	0.204655	
Attrition	-0.239984	1.000000	-0.053892	0.105580	-0.038843	
DailyRate	0.001441	-0.053892	1.000000	-0.009227	-0.015881	
DistanceFromHome	-0.010079	0.105580	-0.009227	1.000000	0.015937	
Education	0.204655	-0.038843	-0.015881	0.015937	1.000000	
EnvironmentSatisfaction	0.008945	-0.101278	0.010620	-0.019730	-0.031925	
HourlyRate	0.034671	-0.036300	0.027128	0.026947	0.017996	
JobInvolvement	0.034193	-0.166036	0.058864	0.010281	0.041046	
JobLevel	0.518333	-0.207634	0.009005	-0.023455	0.093227	
JobSatisfaction	-0.015848	-0.081881	0.032115	-0.004758	-0.003957	
MonthlyIncome	0.511378	-0.193527	0.011030	-0.041201	0.085116	
MonthlyRate	0.025837	0.045744	-0.032211	0.031672	-0.019198	
NumCompaniesWorked	0.296045	0.017279	0.034296	-0.024969	0.126758	
PercentSalaryHike	0.007570	0.002943	0.019325	0.034172	-0.006461	
PerformanceRating	0.005246	0.010728	0.003353	0.020482	-0.020664	
RelationshipSatisfaction	0.058528	-0.020462	0.014539	0.005482	-0.005750	
Shift	0.037117	-0.158322	0.054407	0.029180	0.024451	
TotalWorkingYears	0.692512	-0.234182	0.009378	-0.017663	0.143324	
TrainingTimesLastYear	-0.015408	-0.054836	0.001901	-0.055471	-0.014070	
WorkLifeBalance	-0.004878	-0.090513	-0.028549	-0.037821	0.003933	
YearsAtCompany	0.319012	-0.201373	-0.026892	-0.007420	0.057461	
YearsInCurrentRole	0.222655	-0.207891	0.019651	0.011448	0.051029	
YearsSinceLastPromotion	0.217212	-0.086207	-0.034571	-0.000126	0.045785	
YearsWithCurrManager	0.215909	-0.201087	-0.025272	0.000403	0.055096	

24 rows × 24 columns

dropping rows that exceed a threshold of 0.2

In [ ]:

```
threshold = 0.2
data_df1 = data_df1.drop(data_df1.std()[data_df1.std() < threshold].index.values, axis=1)
```

C:\Users\Joshu\AppData\Local\Temp\ipykernel\_22128\2220865994.py:2: FutureWarning:  
  
Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

In [ ]:

data\_df1

Out[ ]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educa
0	41	0	Travel_Rarely	1102	Cardiology	1	2	Life
1	49	0	Travel_Frequently	279	Maternity	8	1	Life
2	37	1	Travel_Rarely	1373	Maternity	2	2	
3	33	0	Travel_Frequently	1392	Maternity	3	4	Life
4	27	0	Travel_Rarely	591	Maternity	2	1	
...	...	...	...	...	...	...	...	
1671	26	1	Travel_Rarely	471	Neurology	24	3	
1672	46	0	Travel_Rarely	1125	Cardiology	10	3	M
1673	20	0	Travel_Rarely	959	Maternity	1	3	Life
1674	39	0	Travel_Rarely	466	Neurology	1	1	Life
1675	27	0	Travel_Rarely	511	Cardiology	2	2	

1676 rows × 31 columns

In [ ]:

```
#Get dummies
data_df2 = data_df1.copy()
data_df2 = pd.get_dummies(data_df2, drop_first=True)
data_df2
```

Out[ ]:

	Age	Attrition	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate
0	41	0	1102	1	2	2	94
1	49	0	279	8	1	3	61
2	37	1	1373	2	2	4	92
3	33	0	1392	3	4	4	56
4	27	0	591	2	1	1	40
...	...	...	...	...	...	...	...
1671	26	1	471	24	3	3	66
1672	46	0	1125	10	3	3	94
1673	20	0	959	1	3	4	83
1674	39	0	466	1	1	4	65
1675	27	0	511	2	2	1	89

1676 rows × 41 columns

```
In [ ]: data_df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1676 entries, 0 to 1675
Data columns (total 41 columns):
 #   Column                                     Non-Null Count  Dtype
---  -
 0   Age                                       1676 non-null   int64
 1   Attrition                               1676 non-null   int32
 2   DailyRate                               1676 non-null   int64
 3   DistanceFromHome                        1676 non-null   int64
 4   Education                               1676 non-null   int64
 5   EnvironmentSatisfaction                 1676 non-null   int64
 6   HourlyRate                              1676 non-null   int64
 7   JobInvolvement                          1676 non-null   int64
 8   JobLevel                                1676 non-null   int64
 9   JobSatisfaction                         1676 non-null   int64
10   MonthlyIncome                           1676 non-null   int64
11   MonthlyRate                             1676 non-null   int64
12   NumCompaniesWorked                     1676 non-null   int64
13   PercentSalaryHike                      1676 non-null   int64
14   PerformanceRating                      1676 non-null   int64
15   RelationshipSatisfaction                1676 non-null   int64
16   Shift                                   1676 non-null   int64
17   TotalWorkingYears                      1676 non-null   int64
18   TrainingTimesLastYear                  1676 non-null   int64
19   WorkLifeBalance                        1676 non-null   int64
20   YearsAtCompany                         1676 non-null   int64
21   YearsInCurrentRole                     1676 non-null   int64
22   YearsSinceLastPromotion                 1676 non-null   int64
23   YearsWithCurrManager                   1676 non-null   int64
24   BusinessTravel_Travel_Frequently       1676 non-null   uint8
25   BusinessTravel_Travel_Rarely           1676 non-null   uint8
26   Department_Maternity                   1676 non-null   uint8
27   Department_Neurology                   1676 non-null   uint8
28   EducationField_Life Sciences            1676 non-null   uint8
29   EducationField_Marketing                1676 non-null   uint8
30   EducationField_Medical                  1676 non-null   uint8
31   EducationField_Other                    1676 non-null   uint8
32   EducationField_Technical Degree         1676 non-null   uint8
33   Gender_Male                             1676 non-null   uint8
34   JobRole_Administrative                  1676 non-null   uint8
35   JobRole_Nurse                           1676 non-null   uint8
36   JobRole_Other                           1676 non-null   uint8
37   JobRole_Therapist                       1676 non-null   uint8
38   MaritalStatus_Married                   1676 non-null   uint8
39   MaritalStatus_Single                    1676 non-null   uint8
40   OverTime_Yes                            1676 non-null   uint8
dtypes: int32(1), int64(23), uint8(17)
memory usage: 335.7 KB
```

create the final dataframe

```
In [ ]: df_final = data_df2.copy()
df_final
```

Out[ ]:

	Age	Attrition	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate
<b>0</b>	41	0	1102	1	2	2	94
<b>1</b>	49	0	279	8	1	3	61
<b>2</b>	37	1	1373	2	2	4	92
<b>3</b>	33	0	1392	3	4	4	56
<b>4</b>	27	0	591	2	1	1	40
...	...	...	...	...	...	...	...
<b>1671</b>	26	1	471	24	3	3	66
<b>1672</b>	46	0	1125	10	3	3	94
<b>1673</b>	20	0	959	1	3	4	83
<b>1674</b>	39	0	466	1	1	4	65
<b>1675</b>	27	0	511	2	2	1	89

1676 rows × 41 columns

Create a training and test model set

```
In [ ]: from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
```

```
In [ ]: X = df_final.drop(['Attrition'], axis='columns')
```

```
In [ ]: y = df_final['Attrition']
y.to_frame().head()
```

Out[ ]:

	Attrition
<b>0</b>	0
<b>1</b>	0
<b>2</b>	1
<b>3</b>	0
<b>4</b>	0

```
In [ ]: print("X shape",X.shape,"\n","y shape",y.shape)
```

```
X shape (1676, 40)
y shape (1676,)
```

```
In [ ]: # Splitting data
X_train, X_test, y_train, y_test = train_test_split(X, y,
```

```
test_size=.25,  
random_state=42,  
)
```

```
In [ ]: from sklearn.linear_model import LogisticRegression
```

```
In [ ]: logreg = LogisticRegression(penalty='l2', solver='liblinear', max_iter=250)  
logreg.fit(X_train, y_train)
```

```
Out[ ]: LogisticRegression(max_iter=250, solver='liblinear')
```

Predict the model Accuracy

```
In [ ]: y_pred = logreg.predict(X_test)  
print("Model Accuracy:", round(logreg.score(X_test, y_test) * 100, 0), "%")
```

Model Accuracy: 91.0 %